

# PRACTICAL NO.:01

Name: Imam Mahamad Shaikh

Batch:C2

Roll no.:B23025

## Problem Statement:

Predict the price of the Uber ride from a given pickup point to the agreed drop-off location.  
Perform following tasks:

1. Pre-process the dataset.
2. Identify outliers.
3. Check the correlation.
4. Implement linear regression and random forest regression models.
5. Evaluate the models and compare their respective scores like R2, RMSE, etc.

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

df = pd.read_csv('uber.csv')
df
```

	Unnamed: 0	key	fare_amount	\
0	24238194	2015-05-07 19:52:06.0000003	7.5	
1	27835199	2009-07-17 20:04:56.0000002	7.7	
2	44984355	2009-08-24 21:45:00.00000061	12.9	
3	25894730	2009-06-26 08:22:21.0000001	5.3	
4	17610152	2014-08-28 17:47:00.000000188	16.0	
...	...	...	...	...
199995	42598914	2012-10-28 10:49:00.00000053	3.0	
199996	16382965	2014-03-14 01:09:00.0000008	7.5	
199997	27804658	2009-06-29 00:42:00.00000078	30.9	
199998	20259894	2015-05-20 14:56:25.0000004	14.5	
199999	11951496	2010-05-15 04:08:00.00000076	14.1	

```
      pickup_datetime  pickup_longitude  pickup_latitude \
0  2015-05-07 19:52:06 UTC      -73.999817      40.738354
1  2009-07-17 20:04:56 UTC      -73.994355      40.728225
2  2009-08-24 21:45:00 UTC      -74.005043      40.740770
3  2009-06-26 08:22:21 UTC      -73.976124      40.790844
4  2014-08-28 17:47:00 UTC      -73.925023      40.744085
...
199995  2012-10-28 10:49:00 UTC      -73.987042      40.739367
199996  2014-03-14 01:09:00 UTC      -73.984722      40.736837
199997  2009-06-29 00:42:00 UTC      -73.986017      40.756487
199998  2015-05-20 14:56:25 UTC      -73.997124      40.725452
199999  2010-05-15 04:08:00 UTC      -73.984395      40.720077
```

```
      dropoff_longitude  dropoff_latitude  passenger_count
0          -73.999512      40.723217              1
1          -73.994710      40.750325              1
2          -73.962565      40.772647              1
3          -73.965316      40.803349              3
4          -73.973082      40.761247              5
...
199995      -73.986525      40.740297              1
199996      -74.006672      40.739620              1
199997      -73.858957      40.692588              2
199998      -73.983215      40.695415              1
199999      -73.985508      40.768793              1
```

[200000 rows x 9 columns]

```
df.head()
```

```
      Unnamed: 0                 key  fare_amount \
0    24238194  2015-05-07 19:52:06.0000003      7.5
1    27835199  2009-07-17 20:04:56.0000002      7.7
2    44984355  2009-08-24 21:45:00.00000061     12.9
3    25894730  2009-06-26 08:22:21.0000001      5.3
4    17610152  2014-08-28 17:47:00.000000188     16.0
```

```
      pickup_datetime  pickup_longitude  pickup_latitude \
0  2015-05-07 19:52:06 UTC      -73.999817      40.738354
1  2009-07-17 20:04:56 UTC      -73.994355      40.728225
2  2009-08-24 21:45:00 UTC      -74.005043      40.740770
3  2009-06-26 08:22:21 UTC      -73.976124      40.790844
4  2014-08-28 17:47:00 UTC      -73.925023      40.744085
```

```
      dropoff_longitude  dropoff_latitude  passenger_count
0          -73.999512      40.723217              1
1          -73.994710      40.750325              1
2          -73.962565      40.772647              1
3          -73.965316      40.803349              3
4          -73.973082      40.761247              5
```

```
df.tail()
```

```
      Unnamed: 0          key  fare_amount \
199995    42598914  2012-10-28 10:49:00.00000053      3.0
199996    16382965  2014-03-14 01:09:00.0000008       7.5
199997    27804658  2009-06-29 00:42:00.00000078      30.9
199998    20259894  2015-05-20 14:56:25.0000004      14.5
199999    11951496  2010-05-15 04:08:00.00000076      14.1

      pickup_datetime  pickup_longitude  pickup_latitude \
199995  2012-10-28 10:49:00 UTC        -73.987042     40.739367
199996  2014-03-14 01:09:00 UTC        -73.984722     40.736837
199997  2009-06-29 00:42:00 UTC        -73.986017     40.756487
199998  2015-05-20 14:56:25 UTC        -73.997124     40.725452
199999  2010-05-15 04:08:00 UTC        -73.984395     40.720077

      dropoff_longitude  dropoff_latitude  passenger_count
199995        -73.986525      40.740297                  1
199996        -74.006672      40.739620                  1
199997        -73.858957      40.692588                  2
199998        -73.983215      40.695415                  1
199999        -73.985508      40.768793                  1
```

```
df.dtypes
```

```
Unnamed: 0          int64
key              object
fare_amount      float64
pickup_datetime   object
pickup_longitude  float64
pickup_latitude   float64
dropoff_longitude float64
dropoff_latitude  float64
passenger_count    int64
dtype: object
```

```
df.describe()
```

```
      Unnamed: 0  fare_amount  pickup_longitude  pickup_latitude
\count  2.000000e+05  200000.000000  200000.000000  200000.000000
mean    2.771250e+07      11.359955      -72.527638     39.935885
std     1.601382e+07      9.901776      11.437787      7.720539
min     1.000000e+00     -52.000000     -1340.648410     -74.015515
25%    1.382535e+07      6.000000      -73.992065     40.734796
50%    2.774550e+07      8.500000      -73.981823     40.752592
```

```
75%    4.155530e+07      12.500000      -73.967154      40.767158  
max    5.542357e+07      499.000000      57.418457     1644.421482
```

```
dropoff_longitude  dropoff_latitude  passenger_count  
count      199999.000000      199999.000000      200000.000000  
mean       -72.525292       39.923890       1.684535  
std        13.117408       6.794829       1.385997  
min       -3356.666300      -881.985513       0.000000  
25%       -73.991407       40.733823       1.000000  
50%       -73.980093       40.753042       1.000000  
75%       -73.963658       40.768001       2.000000  
max      1153.572603       872.697628      208.000000
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 200000 entries, 0 to 199999  
Data columns (total 9 columns):  
 #   Column           Non-Null Count  Dtype     
---  --  
 0   Unnamed: 0        200000 non-null  int64    
 1   key              200000 non-null  object    
 2   fare_amount      200000 non-null  float64  
 3   pickup_datetime  200000 non-null  object    
 4   pickup_longitude 200000 non-null  float64  
 5   pickup_latitude   200000 non-null  float64  
 6   dropoff_longitude 199999 non-null  float64  
 7   dropoff_latitude  199999 non-null  float64  
 8   passenger_count  200000 non-null  int64    
dtypes: float64(5), int64(2), object(2)  
memory usage: 13.7+ MB
```

```
df.isnull()
```

```
      Unnamed: 0   key   fare_amount  pickup_datetime  
pickup_longitude \\  
0          False  False      False      False  
False  
1          False  False      False      False  
False  
2          False  False      False      False  
False  
3          False  False      False      False  
False  
4          False  False      False      False  
False  
...          ...   ...      ...      ...  
...
```

```
199995      False  False      False      False
False
199996      False  False      False      False
False
199997      False  False      False      False
False
199998      False  False      False      False
False
199999      False  False      False      False
False

          pickup_latitude  dropoff_longitude  dropoff_latitude
passenger_count
0                  False                False      False
False
1                  False                False      False
False
2                  False                False      False
False
3                  False                False      False
False
4                  False                False      False
False
...
...
199995      False                False      False
False
199996      False                False      False
False
199997      False                False      False
False
199998      False                False      False
False
199999      False                False      False
False

[200000 rows x 9 columns]

df.isnull().sum()

Unnamed: 0      0
key              0
fare_amount      0
pickup_datetime  0
pickup_longitude 0
pickup_latitude   0
dropoff_longitude 1
dropoff_latitude  1
passenger_count   0
dtype: int64
```

```

df.shape
(200000, 9)

df.dropna(inplace=True)

df.isnull().sum()

Unnamed: 0      0
key            0
fare_amount    0
pickup_datetime 0
pickup_longitude 0
pickup_latitude  0
dropoff_longitude 0
dropoff_latitude 0
passenger_count 0
dtype: int64

df["pickup_datetime"] = pd.to_datetime(df["pickup_datetime"],
errors="coerce")

df["hour"] = df["pickup_datetime"].dt.hour
df["weekday"] = df["pickup_datetime"].dt.weekday
df["month"] = df["pickup_datetime"].dt.month

# distances in KM
def haversine(lon1, lat1, lon2, lat2):
    lon1, lat1, lon2, lat2 = map(np.radians, [lon1, lat1, lon2, lat2])
    dlon = lon2 - lon1
    dlat = lat2 - lat1
    a = np.sin(dlat/2)**2 + np.cos(lat1) * np.cos(lat2) *
    np.sin(dlon/2)**2
    return 6371 * (2 * np.arcsin(np.sqrt(a))) # km

df["distance"] = haversine(df["pickup_longitude"],
df["pickup_latitude"],
                           df["dropoff_longitude"],
df["dropoff_latitude"])

df

      Unnamed: 0      key  fare_amount \
0      24238194  2015-05-07 19:52:06.0000003   7.5
1      27835199  2009-07-17 20:04:56.0000002   7.7
2      44984355  2009-08-24 21:45:00.00000061  12.9
3      25894730  2009-06-26 08:22:21.0000001   5.3
4      17610152  2014-08-28 17:47:00.000000188  16.0
..        ...
199995    42598914  2012-10-28 10:49:00.00000053   3.0
199996    16382965  2014-03-14 01:09:00.0000008   7.5

```

199997	27804658	2009-06-29	00:42:00	0.00000078	30.9
199998	20259894	2015-05-20	14:56:25	0.0000004	14.5
199999	11951496	2010-05-15	04:08:00	0.00000076	14.1
<hr/>					
0	2015-05-07	19:52:06+00:00	-73.999817	40.738354	
1	2009-07-17	20:04:56+00:00	-73.994355	40.728225	
2	2009-08-24	21:45:00+00:00	-74.005043	40.740770	
3	2009-06-26	08:22:21+00:00	-73.976124	40.790844	
4	2014-08-28	17:47:00+00:00	-73.925023	40.744085	
<hr/>					
199995	2012-10-28	10:49:00+00:00	-73.987042	40.739367	
199996	2014-03-14	01:09:00+00:00	-73.984722	40.736837	
199997	2009-06-29	00:42:00+00:00	-73.986017	40.756487	
199998	2015-05-20	14:56:25+00:00	-73.997124	40.725452	
199999	2010-05-15	04:08:00+00:00	-73.984395	40.720077	
<hr/>					
weekday		dropoff_longitude	dropoff_latitude	passenger_count	hour
0		-73.999512	40.723217		1 19
3		-73.994710	40.750325		1 20
1		-73.962565	40.772647		1 21
4		-73.965316	40.803349		3 8
0		-73.973082	40.761247		5 17
3		<hr/>			
<hr/>					
199995		-73.986525	40.740297		1 10
6		-74.006672	40.739620		1 1
199996		-73.858957	40.692588		2 0
4		-73.983215	40.695415		1 14
0		-73.985508	40.768793		1 4
2		<hr/>			
199999		<hr/>			
5		<hr/>			
<hr/>					
0	month	distance			
0	5	1.683323			
1	7	2.457590			
2	8	5.036377			
3	6	1.661683			
4	8	4.475450			
<hr/>					
199995	10	0.112210			

```

199996      3   1.875050
199997      6  12.850319
199998      5   3.539715
199999      5   5.417783

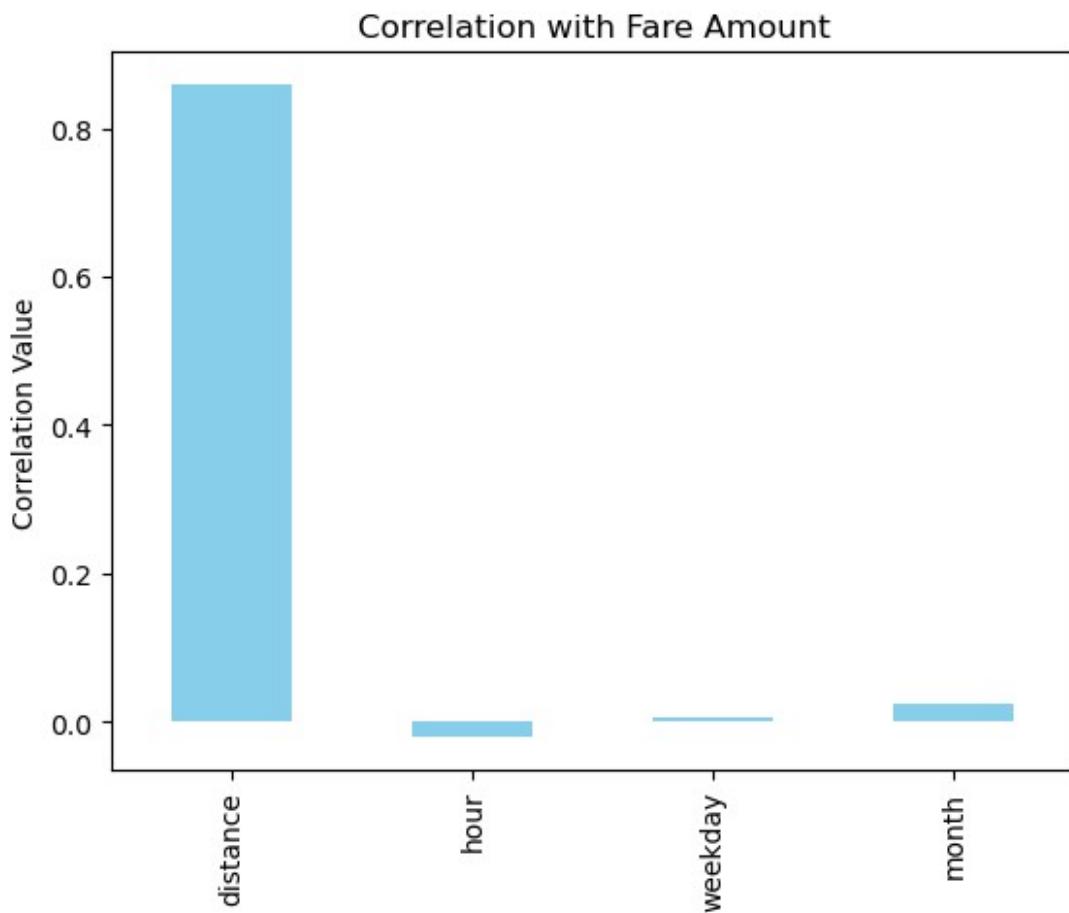
[199999 rows x 13 columns]

# Outlier handling
df = df[(df["fare_amount"] > 0) & (df["fare_amount"] < 100)]
df = df[(df["distance"] > 0) & (df["distance"] < 100)]

corr = df[["fare_amount", "distance", "hour", "weekday", "month"]].corr()
["fare_amount"]

corr.drop("fare_amount").plot(kind="bar", color="skyblue")
plt.title("Correlation with Fare Amount")
plt.ylabel("Correlation Value")
plt.show()

```



```

# 5. Features & Target
X = df[["distance", "hour", "weekday", "month", "passenger_count"]]
y = df["fare_amount"]

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor

#LinearRegression
lr = LinearRegression()
lr.fit(X_train, y_train)
pred_lr = lr.predict(X_test)

pred_lr

array([ 5.80204908,  8.40852548,  8.54483812, ...,  9.47122511,
       26.02043998,  5.69969744])

rf = RandomForestRegressor(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)
pred_rf = rf.predict(X_test)

pred_rf

array([ 6.043,   9.023,   9.299, ...,  9.286,  26.837,   6.523])

from sklearn.metrics import r2_score, mean_squared_error,
mean_absolute_error

results = []
def evaluate(y_test, y_pred, name):
    r2 = r2_score(y_test, y_pred)
    rmse = np.sqrt(mean_squared_error(y_test, y_pred))
    mae = mean_absolute_error(y_test, y_pred)
    results.append([name, r2, rmse, mae])
    print(f"{name} -> R2: {r2:.3f}, RMSE: {rmse:.2f}, MAE: {mae:.2f}")

evaluate(y_test, pred_lr, "Linear Regression")
Linear Regression -> R2: 0.712, RMSE: 5.05, MAE: 2.38

evaluate(y_test, pred_rf, "Random Forest")
Random Forest -> R2: 0.787, RMSE: 4.34, MAE: 2.39

results = np.array(results, dtype=object)

# Plot bar chart for comparison
metrics = ["R2", "RMSE", "MAE"]

```

```

x = np.arange(len(metrics)) # [0,1,2]
width = 0.35

plt.figure(figsize=(8,5))
plt.bar(x - width/2, results[0,1:].astype(float), width,
label=results[0,0])
plt.bar(x + width/2, results[1,1:].astype(float), width,
label=results[1,0])

plt.xticks(x, metrics)
plt.ylabel("Score / Error")
plt.title("Model Performance Comparison")
plt.legend()
plt.show()

```

