

PRACTICAL NO.:01

Name: Imam Mahamad Shaikh

Batch:C2

Roll no.:B23025

Problem Statement:

Given a bank customer, build a neural network-based classifier that can determine whether they will leave or not in the next 6 months. Dataset Description: The case study is from an open-source dataset from Kaggle. The dataset contains 10,000 sample points with 14 distinct features such as CustomerId, CreditScore, Geography, Gender, Age, Tenure, Balance, etc. Link to the Kaggle project: <https://www.kaggle.com/barelydedicated/bank-customer-churn-modeling>
Perform following steps:

1. Read the dataset.
2. Distinguish the feature and target set and divide the data set into training and test sets.
3. Normalize the train and test data.
4. Initialize and build the model. Identify the points of improvement and implement the same.
5. Print the accuracy score and confusion matrix (5 points).

```
import pandas as pd
import numpy as np

df = pd.read_csv("Churn_Modelling.csv")

df
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender
Age \						
0	1	15634602	Hargrave	619	France	Female
42	2	15647311	Hill	608	Spain	Female
41	3	15619304	Onio	502	France	Female

42		4	15701354	Boni	699	France	Female	
39		5	15737888	Mitchell	850	Spain	Female	
43		
...		
9995	9996	15606229	Obijaku	771	France	Male		
39		9997	15569892	Johnstone	516	France	Male	
35		9998	15584532	Liu	709	France	Female	
36		9999	15682355	Sabbatini	772	Germany	Male	
42		10000	15628319	Walker	792	France	Female	
28								
	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember		\	
0	2	0.00	1	1	1		1	
1	1	83807.86	1	0	1			
2	8	159660.80	3	1	0		0	
3	1	0.00	2	0	0		0	
4	2	125510.82	1	1	1		1	
...	
9995	5	0.00	2	1	0		0	
9996	10	57369.61	1	1	1		1	
9997	7	0.00	1	0	1			
9998	3	75075.31	2	1	0		0	
9999	4	130142.79	1	1	0			
	EstimatedSalary	Exited						
0	101348.88	1						
1	112542.58	0						
2	113931.57	1						
3	93826.63	0						
4	79084.10	0						
...						
9995	96270.64	0						
9996	101699.77	0						
9997	42085.58	1						
9998	92888.52	1						
9999	38190.78	0						
[10000 rows x 14 columns]								
df.head()								
	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	\

0	1	15634602	Hargrave	619	France	Female	42
1	2	15647311	Hill	608	Spain	Female	41
2	3	15619304	Onio	502	France	Female	42
3	4	15701354	Boni	699	France	Female	39
4	5	15737888	Mitchell	850	Spain	Female	43
0	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember		
0	2	0.00	1	1	1	1	1
1	1	83807.86	1	0	0	1	1
2	8	159660.80	3	1	0	0	0
3	1	0.00	2	0	0	0	0
4	2	125510.82	1	1	1	1	1
0	EstimatedSalary	Exited					
0	101348.88	1					
1	112542.58	0					
2	113931.57	1					
3	93826.63	0					
4	79084.10	0					
df.tail()							
Age \	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	
9995	9996	15606229	Obijiaku	771	France	Male	
39							
9996	9997	15569892	Johnstone	516	France	Male	
35							
9997	9998	15584532	Liu	709	France	Female	
36							
9998	9999	15682355	Sabbatini	772	Germany	Male	
42							
9999	10000	15628319	Walker	792	France	Female	
28							
9995	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember		
9995	5	0.00	2	1	0	1	0
9996	10	57369.61	1	1	1	1	1
9997	7	0.00	1	0	1	0	1
9998	3	75075.31	2	1	0	0	0
9999	4	130142.79	1	1	1	0	0
9995	EstimatedSalary	Exited					
9995	96270.64	0					
9996	101699.77	0					
9997	42085.58	1					

```

9998      92888.52      1
9999      38190.78      0

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   RowNumber        10000 non-null   int64  
 1   CustomerId       10000 non-null   int64  
 2   Surname          10000 non-null   object  
 3   CreditScore      10000 non-null   int64  
 4   Geography         10000 non-null   object  
 5   Gender            10000 non-null   object  
 6   Age               10000 non-null   int64  
 7   Tenure            10000 non-null   int64  
 8   Balance           10000 non-null   float64 
 9   NumOfProducts     10000 non-null   int64  
 10  HasCrCard        10000 non-null   int64  
 11  IsActiveMember    10000 non-null   int64  
 12  EstimatedSalary   10000 non-null   float64 
 13  Exited            10000 non-null   int64  
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB

df.describe()

      RowNumber  CustomerId  CreditScore      Age
Tenure \
count  10000.00000  1.000000e+04  10000.00000  10000.00000
10000.00000
mean   5000.50000  1.569094e+07  650.528800  38.921800
5.012800
std    2886.89568  7.193619e+04  96.653299  10.487806
2.892174
min    1.00000  1.556570e+07  350.000000  18.000000
0.000000
25%   2500.75000  1.562853e+07  584.000000  32.000000
3.000000
50%   5000.50000  1.569074e+07  652.000000  37.000000
5.000000
75%   7500.25000  1.575323e+07  718.000000  44.000000
7.000000
max   10000.00000  1.581569e+07  850.000000  92.000000
10.000000

      Balance  NumOfProducts  HasCrCard  IsActiveMember \
count  10000.00000  10000.00000  10000.00000  10000.00000

```

```
mean    76485.889288      1.530200      0.70550      0.515100
std     62397.405202      0.581654      0.45584      0.499797
min     0.000000      1.000000      0.00000      0.000000
25%    0.000000      1.000000      0.00000      0.000000
50%    97198.540000      1.000000      1.00000      1.000000
75%   127644.240000      2.000000      1.00000      1.000000
max   250898.090000      4.000000      1.00000      1.000000
```

```
EstimatedSalary      Exited
count    10000.000000  10000.000000
mean     100090.239881  0.203700
std      57510.492818  0.402769
min     11.580000      0.000000
25%    51002.110000      0.000000
50%    100193.915000      0.000000
75%    149388.247500      0.000000
max   199992.480000      1.000000
```

```
df.dtypes
```

```
RowNumber          int64
CustomerId         int64
Surname            object
CreditScore        int64
Geography          object
Gender              object
Age                int64
Tenure              int64
Balance             float64
NumOfProducts       int64
HasCrCard           int64
IsActiveMember      int64
EstimatedSalary     float64
Exited              int64
dtype: object
```

```
df.isnull()
```

```
      RowNumber CustomerId Surname CreditScore Geography Gender
Age \
0      False      False    False     False    False    False
1      False      False    False     False    False    False
2      False      False    False     False    False    False
3      False      False    False     False    False    False
4      False      False    False     False    False    False
```

...
...
9995	False	False	False	False	False	False	False
False							
9996	False	False	False	False	False	False	False
False							
9997	False	False	False	False	False	False	False
False							
9998	False	False	False	False	False	False	False
False							
9999	False	False	False	False	False	False	False
False							
0	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\\	
1	False	False	False	False	False	False	
2	False	False	False	False	False	False	
3	False	False	False	False	False	False	
4	False	False	False	False	False	False	
...	
9995	False	False	False	False	False	False	
9996	False	False	False	False	False	False	
9997	False	False	False	False	False	False	
9998	False	False	False	False	False	False	
9999	False	False	False	False	False	False	
0	EstimatedSalary	Exited					
1	False	False					
2	False	False					
3	False	False					
4	False	False					
...					
9995	False	False					
9996	False	False					
9997	False	False					
9998	False	False					
9999	False	False					

[10000 rows x 14 columns]

df.isnull().sum()

RowNumber	0
CustomerId	0
Surname	0
CreditScore	0
Geography	0
Gender	0
Age	0

```

Tenure          0
Balance         0
NumOfProducts   0
HasCrCard       0
IsActiveMember  0
EstimatedSalary 0
Exited          0
dtype: int64

x = df.drop(columns=["RowNumber", "CustomerId", "Surname", "Exited"])
y = df["Exited"]

x

      CreditScore Geography Gender Age Tenure Balance
NumOfProducts \
0            619     France Female  42      2    0.00
1
1            608     Spain  Female  41      1  83807.86
1
2            502     France Female  42      8 159660.80
3
3            699     France Female  39      1    0.00
2
4            850     Spain  Female  43      2 125510.82
1
...
...
9995          771     France Male   39      5    0.00
2
9996          516     France Male   35     10  57369.61
1
9997          709     France Female  36      7    0.00
1
9998          772     Germany Male   42      3  75075.31
2
9999          792     France Female  28      4 130142.79
1

      HasCrCard IsActiveMember EstimatedSalary
0            1             1        101348.88
1            0             1        112542.58
2            1             0        113931.57
3            0             0        93826.63
4            1             1        79084.10
...
9995          1             0        96270.64
9996          1             1        101699.77
9997          0             1        42085.58
9998          1             0        92888.52

```

```
9999           1           0        38190.78
[10000 rows x 10 columns]
y
0      1
1      0
2      1
3      0
4      0
...
9995    0
9996    0
9997    1
9998    1
9999    0
Name: Exited, Length: 10000, dtype: int64
from sklearn.preprocessing import OneHotEncoder, StandardScaler
X = pd.get_dummies(x, drop_first=True)
X=X.values
X
array([[619, 42, 2, ..., False, False, False],
       [608, 41, 1, ..., False, True, False],
       [502, 42, 8, ..., False, False, False],
       ...,
       [709, 36, 7, ..., False, False, False],
       [772, 42, 3, ..., True, False, True],
       [792, 28, 4, ..., False, False, False]], dtype=object)

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
from sklearn.neural_network import MLPClassifier
mlp = MLPClassifier(hidden_layer_sizes=(32, 16), max_iter=200,
random_state=42)
mlp.fit(X_train, y_train)
C:\ProgramData\anaconda3\Lib\site-packages\sklearn\neural_network\
_multilayer_perceptron.py:690: ConvergenceWarning: Stochastic
```

```
Optimizer: Maximum iterations (200) reached and the optimization
hasn't converged yet.
warnings.warn(
    "Maximum iterations (200) reached and the optimization hasn't
    converged yet. Try increasing max_iter if necessary."
)

MLPClassifier(hidden_layer_sizes=(32, 16), random_state=42)
y_pred = mlp.predict(X_test)
y_pred
array([0, 0, 0, ..., 1, 0, 0], dtype=int64)
from sklearn.metrics import accuracy_score, confusion_matrix
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))

Accuracy: 0.861
Confusion Matrix:
[[1519  88]
 [ 190 203]]
```